

PyAMG

Algebraic Multigrid Solvers in Python

PyAMG

Algebraic Multigrid Solvers in Python

Nathan Bell, Nvidia

Luke Olson, University of Illinois

Jacob Schroder, Lawrence Livermore National Lab

Copper Mountain 2012

Boot disc

- ❖ For Mac
 - ❖ Press and hold `c` immediately after reboot
- ❖ For the usual suspects (Dell, HP, Thinkpad, etc...)
 - ❖ Some automatically detect a bootable DVD
 - ❖ Pressing `F12`, `F8`, or `F6` at boot window can list boot options
 - ❖ Last, an advanced approach changes BIOS boot order
 - ❖ Press `F1`, `delete`, or `esc` during initial boot window
 - ❖ From the BIOS, change boot order to start with DVD drive first
- ❖ After booting, open a terminal (Applications—>Accessories—>Terminal)

Task 0.1: Installing PyAMG

Compile PyAMG (if using boot disc)

```
$ tar -xvf pyamg2.0.1.tar.gz  
$ cd pyamg/  
$ sudo python setup.py install  
$ cd Examples/WorkshopCopper12  
$ ipython
```

Test PyAMG during interactive iPython session, enter:

```
import pyamg  
pyamg.test()
```

What is PyAMG

- ✧ Algebraic multigrid (AMG) workbench
 - ✧ Readable and reproducible AMG
 - ✧ Efficient serial solver
- ✧ Python-based
 - ✧ Readability and useability
 - ✧ C++ backend for speed
- ✧ BLAS, LAPACK, optimized routines, etc...

Goal 1: Ease-of-use

- ❖ Accessible interface to non-experts
- ❖ Extensive documentation and references
- ❖ Portability through modular multiplatform Python libraries
- ❖ Rapid prototyping of new techniques
 - ❖ Organize source code into intuitive reusable components
 - ❖ High-level Python allows for rapid swapping of components

Goal 2: Speed

- ❖ Solve millions of unknowns on laptop / desktop
- ❖ Use hybrid coding strategy (Python as glue)
 - ❖ Performance sensitive portions are small fraction of code
 - ❖ 80% Python / 20% natively compiled C/C++/Fortran
- ❖ Example:
 - ❖ High-level multigrid cycling in Python
 - ❖ Calls `gauss_seidel(A,x,b,iterations=1)`
 - ❖ All computation done in C++ routine

PyAMG features

- ❖ Ruge-Stüben AMG
- ❖ Smoothed aggregation (standard and adaptive)
- ❖ Native complex support
- ❖ Nonsymmetric matrices
- ❖ Krylov solvers (CG, BiCGStab, GMRES, fGMRES, CGNR, CR)
- ❖ Compatible relaxation (experimental)
- ❖ Relaxation methods (GS, wJ, SOR, Kacz, Cheby, Schwarz)
- ❖ Visualizations (Paraview, Matplotlib)

Dependencies

PyAMG

Multilevel solvers, relaxation methods, Krylov methods

Scipy

*LAPACK and
sparse matrix
operations*

C++/Swig

*Easy interface from
Python to C++*

Nose

*Unit tests, i.e.,
does everything
work?*

Matplotlib

Visualizations

iPython

*Python interpreter,
interactive sessions*

Numpy

*Array operations
(BLAS)*

What PyAMG does not do...

- ❖ Solve everything in linear time
- ❖ Multicore (coming)
- ❖ GPU acceleration (Cusp)
- ❖ Large-scale parallel simulations
 - ❖ See Hypre (Livermore) and ML (Sandia)

General Structure

- ❖ `multilevel.py` (multilevel solver class)
- ❖ Main structure for hierarchy (SA or RS)
- ❖ Handles cycling and coarse solver
- ❖ Contains list of level object instances
- ❖ Each level object instance contains:
 - ❖ Matrices: A , P , R
 - ❖ Functions: `presmoothen`, `postsmoothen`

Starting from pyamg/pyamg

- * `multilevel.py` Multilevel solver class
- * `strength.py` Strength-of-connection routines
- * `classical/` Ruge-Stüben construction routines
- * `aggregation/` SA construction routines
- * `krylov/` Krylov solvers, e.g., CG, GMRES, fGMRES
- * `relaxation/` Relaxation methods, e.g., GS, wJ, Cheby
- * `Graph/` Graph algorithms for coarsening, e.g., MIS
- * `vis/` Visualizations in Matplotlib and Paraview
- * `amg_core/` C++ functions called through Swig
- * `gallery/` Construct model problems
- * `util/` Utility functions, e.g., spectral radius

Task 1.1: Getting used to Python

- ❖ Accessing documentation inside of iPython
 - ❖ Use <tab> on object to see its members
 - ❖ Use ? on object or function for documentation
 - ❖ Use spacebar to page, and q to quit documentation screen
- ❖ Inside of iPython (\$ ipython), enter:

*The most useful things
you'll learn today*

```
from pyamg import gallery, smoothed_aggregation_solver
```

```
A = gallery.poisson( (50,50), format='csr')
```

```
ml = smoothed_aggregation_solver(A)
```

```
ml.<tab>
```

```
ml.__class__
```

```
ml.__doc__
```

```
ml.__init__
```

```
ml.__module__
```

```
ml.__repr__
```

```
ml._multilevel_solver__solve
```

```
ml.aspreconditioner
```

```
ml.coarse_solver
```

```
ml.cycle_complexity
```

```
ml.grid_complexity
```

```
ml.level
```

```
ml.levels
```

```
ml.operator_complexity
```

```
ml.psolve
```

```
ml.solve
```

```
ml.solve?
```

```
gallery.poisson?
```

Task 1.2: Sparse Matrices

- ❖ Construct a sparse matrix
- ❖ Inside of iPython, enter:

*If you ever get lost, exit iPython
(ctrl-d, ctrl-d), re-enter iPython,
and type `run taskx.x.py`*

```
from scipy.sparse import *  
csr_matrix?  
from numpy import array  
row = array([0,0,1,2,2,2])  
col = array([0,2,2,0,1,2])  
data = array([1,2,3,4,5,6])  
B = csr_matrix( (data,(row,col)), shape=(3,3) )  
B.<tab>  
print(B.todense())  
    [[1 0 2]  
     [0 0 3]  
     [4 5 6]]  
B = B.tocoo()
```

Task 1.3: Using the gallery

- ❖ Generate a sparse matrix by running script
- ❖ Inside of iPython, enter:

```
run task1.3

print(A[5050,:].data)
    [-0.22 -0.25  0.22 -0.75  2.   -0.75  0.22 -0.25 -0.22]

print(sten)
    [[-0.22 -0.25  0.22]
     [-0.75  2.   -0.75]
     [ 0.22 -0.25 -0.22]]
```

Script: task1.3.py

```
from pyamg.gallery.diffusion import diffusion_stencil_2d
from pyamg.gallery import stencil_grid
from numpy import set_printoptions
set_printoptions(precision=2)
sten = diffusion_stencil_2d(type='FD', \
    epsilon=0.001, theta=3.1416/3.0)
A = stencil_grid(sten, (100,100), format='csr')
```


Task 1.4: Building a MG hierarchy

- ✧ Inside of iPython, enter:

```
from pyamg import *
ml = smoothed_aggregation_solver(A)
print(ml)
    multilevel_solver
    Number of Levels:      3
    Operator Complexity:   1.126
    Grid Complexity:       1.130
    Coarse Solver:         'pinv2'
      level  unknowns      nonzeros
        0      10000      88804 [88.84%]
        1       1156      10000 [10.00%]
        2        144       1156 [ 1.16%]

print(ml.levels[0].A.shape)
(10000, 10000)

print(ml.levels[0].P.shape)
(10000, 1156)

print(ml.levels[0].R.shape)
(1156, 10000)
```

Task 1.5: Solving a problem

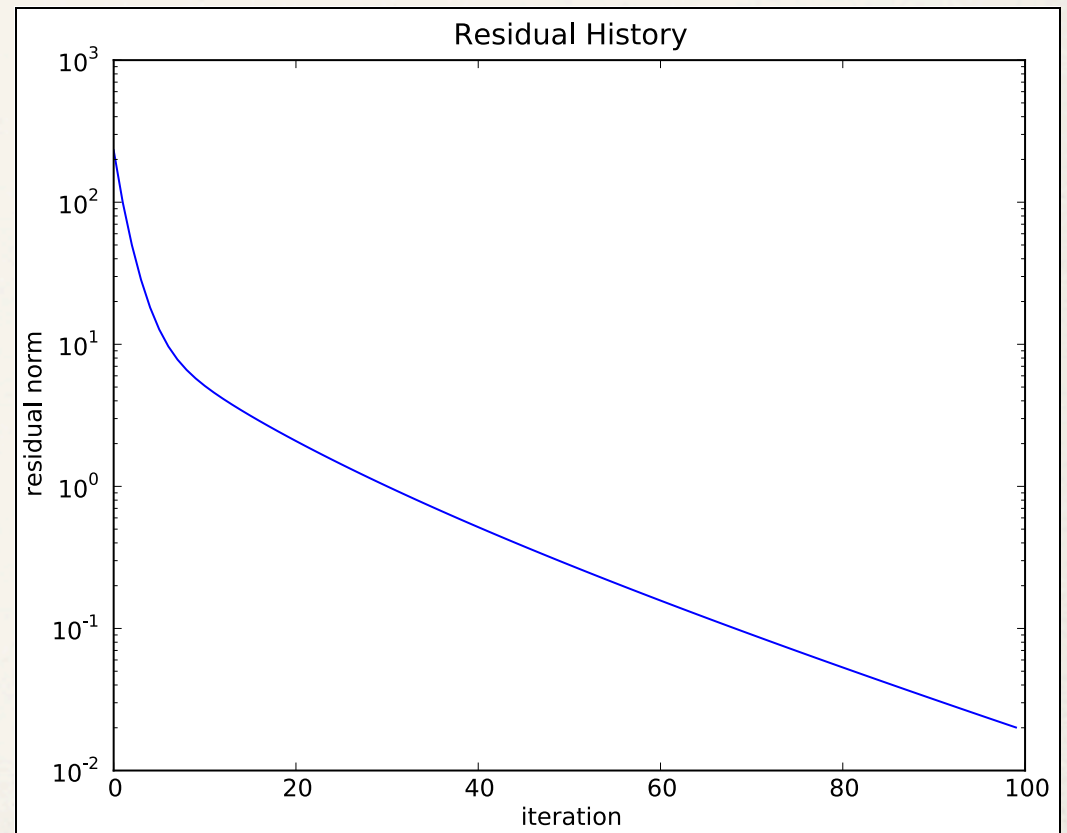
- ❖ Inside of iPython, enter:

```
run task1.5
```

Script: task1.5.py

```
from numpy import ones
b = ones((A.shape[0],1))
res = []
x = ml.solve(b, tol=1e-8, \
             residuals=res)

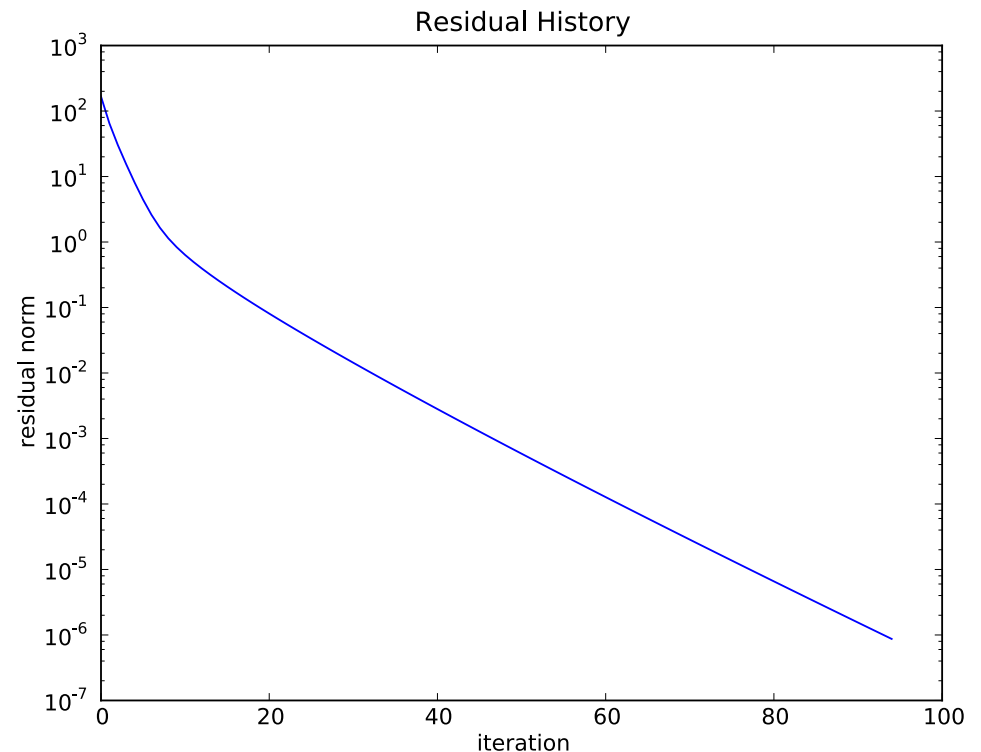
from pylab import *
semilogy(res[1:])
xlabel('iteration')
ylabel('residual norm')
title('Residual History')
show()
```



Task 1.6: Changing MG options

- ❖ Use advanced coarsening and prolongation smoothing options
- ❖ Inside of iPython, enter:

```
from pyamg import *
from numpy import ones
ml = smoothed_aggregation_solver(A, \
    strength='evolution', \
    smooth=('energy', {'degree':4}))
b = ones((A.shape[0],1))
res = []
x = ml.solve(b, tol=1e-8, residuals=res)
from pylab import *
semilogy(res[1:])
xlabel('iteration')
ylabel('residual norm')
title('Residual History')
show()
```



Intermediate Tasks

- ❖ Modify existing multilevel hierarchy
- ❖ Add new prolongation smoothing function to PyAMG source
- ❖ Visualizations with Paraview
- ❖ Loading matrix from file
- ❖ Blackbox solve

Task 2.1: Modifying the hierarchy

- ❖ Modify existing multilevel solver object
- ❖ Replace existing pre/post-smoothers with new user-provided routine
- ❖ Execute commands in shell:

ctrl-d, ctrl-d (exit iPython)

\$ python task2.1.py

Set new pre/post-smoother {

```
def new_relax(A,x,b):  
    x[:] += 0.125*(b - A*x)  
  
A = gallery.poisson( (100,100), format='csr')  
b = ones( (A.shape[0],1))  
res = []  
ml = smoothed_aggregation_solver(A)  
ml.levels[0].presmoothing = new_relax  
ml.levels[0].postsolving = new_relax  
x = ml.solve(b, tol=1e-8, residuals=res)  
  
semilogy(res[1:])  
show()
```

Task 2.2: Adding a smoother

\$ gedit ~/pyamg/pyamg/aggregation/aggregation.py
At line 409, insert two new lines:

New
Lines {

```
if fn == 'jacobi':  
    P = jacobi_prolongation_smoother(A, T, C, B, **kwargs)  
elif fn == 'simple':  
    P = T - 0.2*A*T  
elif fn == 'richardson':  
    ...
```

```
$ cd ~/pyamg/  
$ sudo python setup.py install  
$ cd Examples/WorkshopCopper12  
$ python task2.2.py
```

```
A = gallery.poisson( (100,100), format='csr')  
ml = smoothed_aggregation_solver(A, smooth='simple')  
...
```


Task 2.3: Plotting aggregates

❖ Visualization with Paraview

```
$ python task2.3.py
```

```
$ paraview&
```

load data

```
{ data = load_example('unit_square')  
  A = data['A'].tocsr()  
  V = data['vertices']  
  E2V = data['elements']
```

create hierarchy

```
{ ml = smoothed_aggregation_solver(A, keep=True, max_coarse=10)  
  b = sin(pi*V[:,0])*sin(pi*V[:,1])  
  x = ml.solve(b)
```

save aggregates

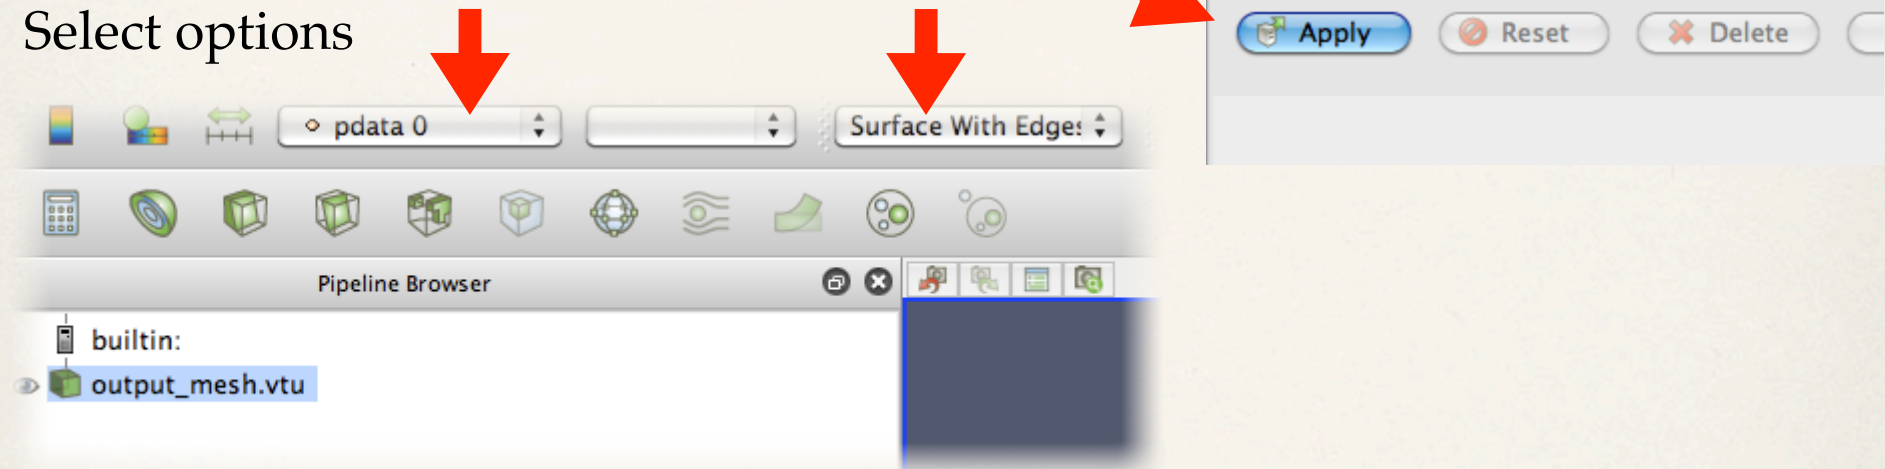
```
{ vis_coarse.vis_aggregate_groups(Verts=V, E2V=E2V,  
  Agg=ml.levels[0].AggOp, mesh_type='tri',  
  output='vtk', fname='output_aggs.vtu')
```

save mesh

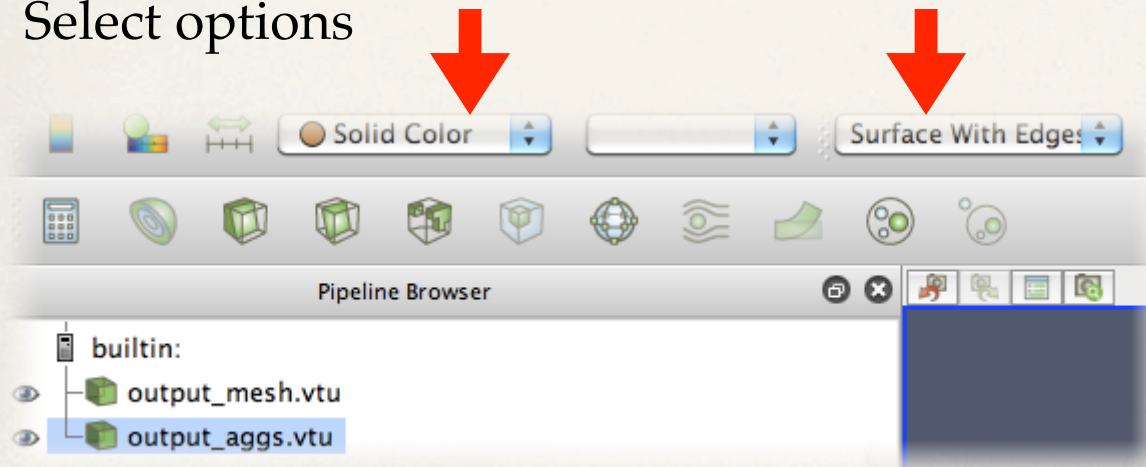
```
{ vtk_writer.write_basic_mesh(Verts=V, E2V=E2V,  
  pdata = x,  
  mesh_type='tri',  
  fname='output_mesh.vtu')
```

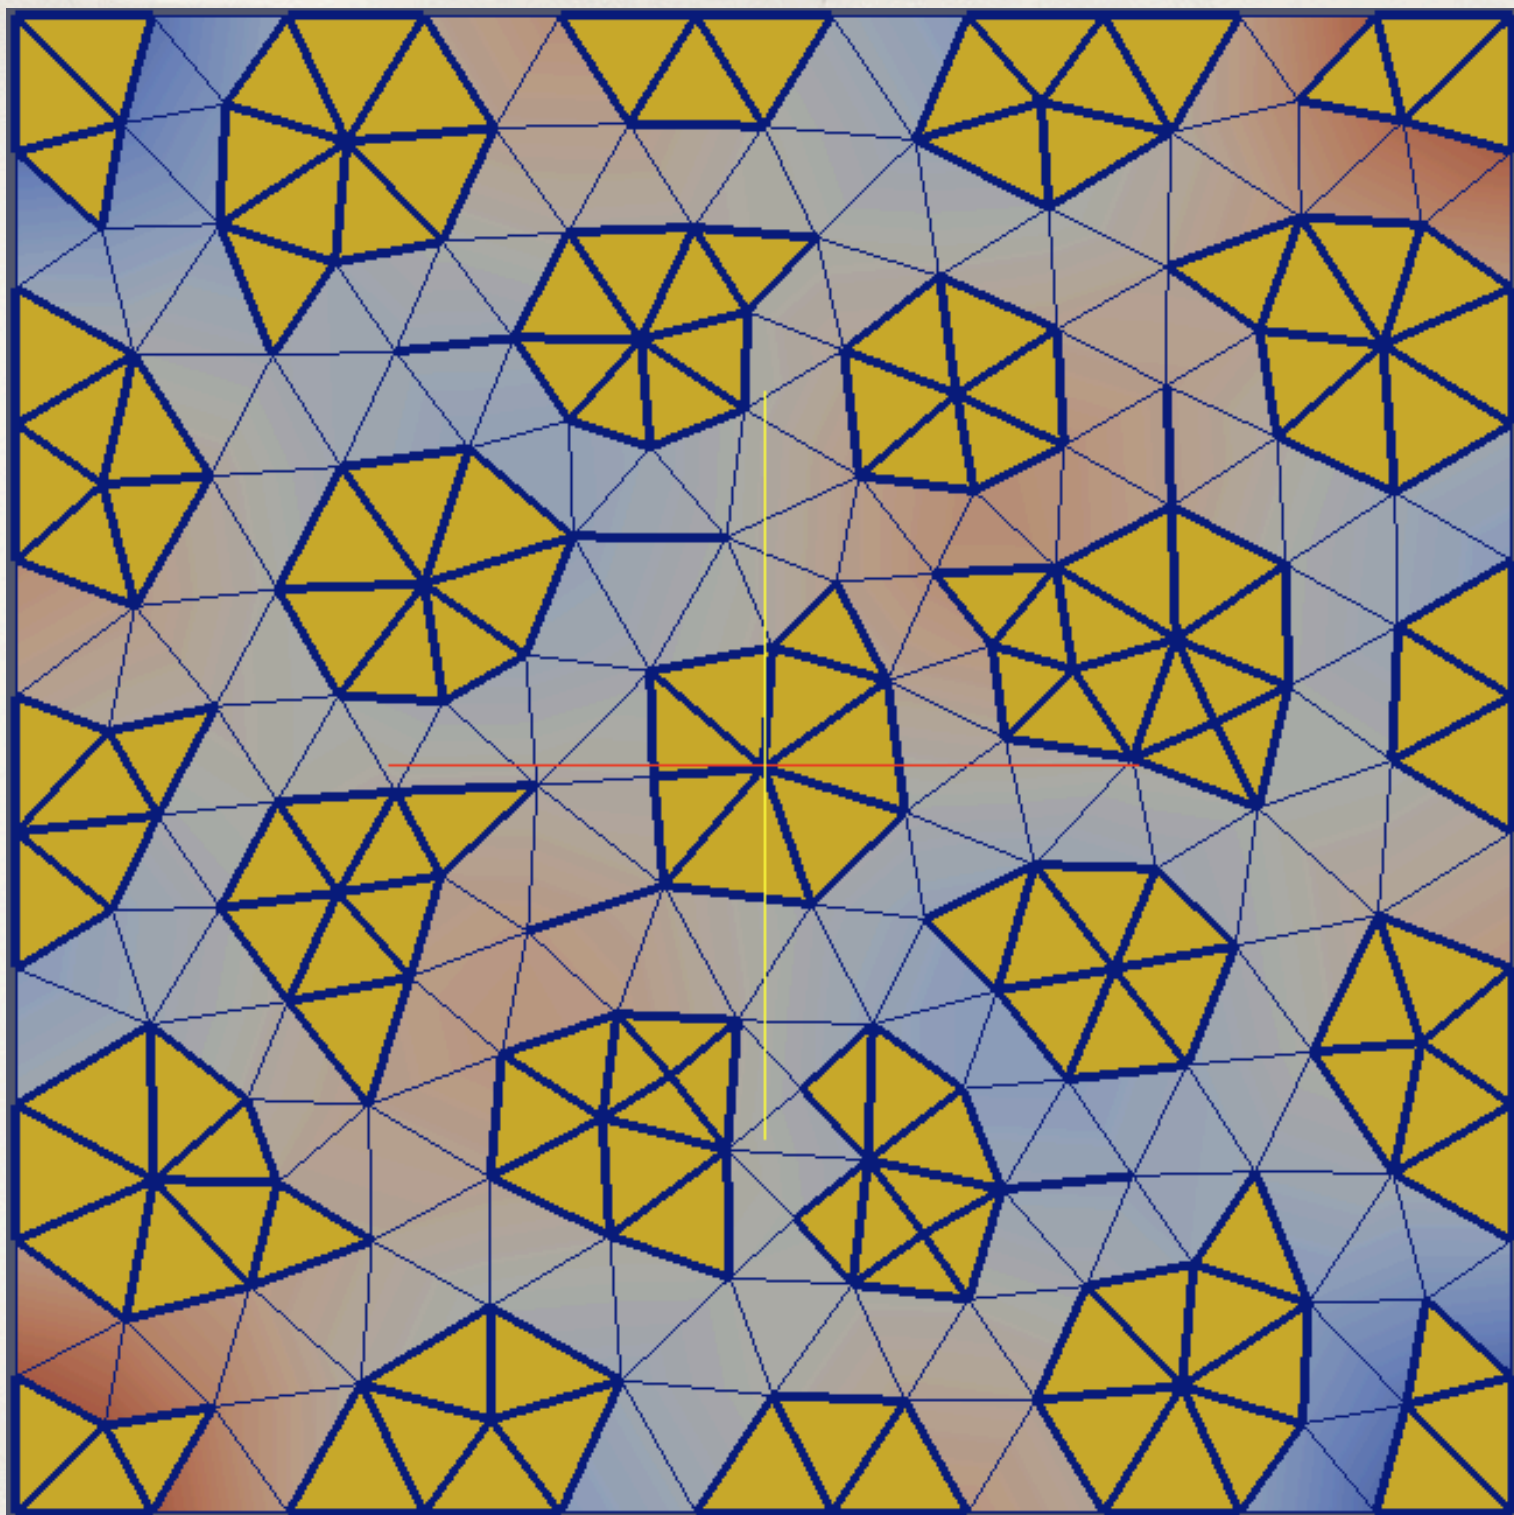
Task 2.3: Plot mesh and aggregates

- ❖ Open `output_mesh.vtu`, then click apply
- ❖ Select options



- ❖ Open `output_aggs.vtu`, then click apply
- ❖ Select options





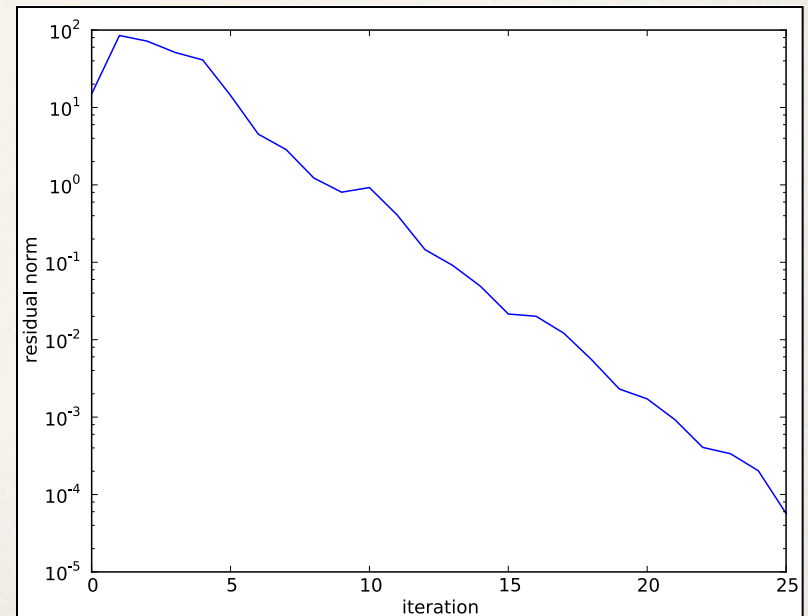
Task 2.4: Loading a matrix

nonsymmetric

```
$ python task2.4.py
```

```
data = loadmat('../..//pyamg/gallery/example_data/recirc_flow.mat')
A = data['A']

from pyamg import *
ml = smoothed_aggregation_solver(A, symmetry='nonsymmetric', max_coarse=5)
```



Try it with...

```
ml = smoothed_aggregation_solver(A, symmetry='symmetric', max_coarse=5)
```

Task 2.5: Running blackbox solve

Inside of iPython, enter:

- ❖ “blackbox” solve attempts to pick the most robust options
- ❖ solve once
- ❖ solve again (same hierarchy)



```
from scipy import rand
from numpy import arange, array
from pyamg import solve
from pyamg.gallery import poisson
from pyamg.util.linalg import norm

# Run solve(...) with the verbose option
n = 100
A = poisson((n,n),format='csr')
b = array(arange(A.shape[0]))
x = solve(A,b,verb=True)

# Return the solver for re-use
(x,ml) = solve(A, b, verb=True,
               return_solver=True, tol=1e-8)

# Run for a new right-hand-side
b2 = rand(b.shape[0],)
x2 = solve(A, b2, verb=True,
           existing_solver=ml, tol=1e-8)
```

Advanced Tasks

- ❖ SWIG

- ❖ SWIG interfaces between C++ and Python
- ❖ Replace slow Python segments with C++
- ❖ This task compares pure Python and hybrid Python/C++ versions of forward and backward substitution

- ❖ Important files for example

- ❖ `numpy.i` interface between NumPy and C++ (esp. arrays)
- ❖ `complex_ops.h` interface for complex data types
- ❖ `splinalg.i` IN, INPLACE and OUT types, templating
- ❖ `splinalg.h` headers, function definitions (plain vanilla C++)

```
$ cd ~/pyamg/Examples/SWIG
```


Task 3.1: Calling C++

splinalg.h defines the C++:

```
template<class I, class T>
void forwardsolve(const I Ap[], const I Aj[], const T Ax[],
                  T x[], const T b[], const I n)
{...
}
```

splinalg.i defines the C++ interface for SWIG:

```
/* INPLACE types */
#define T_INPLACE_ARRAY1( ctype )
%apply ctype * INPLACE_ARRAY {
    ctype    x [ ]
};
```

SWIG compiles and creates the python interface:

```
$ swig -c++ -python splinalg.i
```

```
splinalg.forwardsolve(L.indptr,L.indices,L.data,x,b,n)
```

Task 3.1: Calling C++

To compile example:

```
ctrl-d, ctrl-d (exit iPython)  
$ cd ~/pyamg/Examples/SWIG  
$ sudo python setup.py install
```

Run example calling C++ routines with SWIG:

```
$ python testbasic.py
```

Task 3.1: Calling C++

- ❖ C++ call:

time for one LU solve = 0.1998 ms

- ❖ \$ gedit precondition.py

- ❖ change line 74 to

{

```
def preconditioner_matvec(L,U):  
    def matvec(x):  
        return lusolve_reference(L,U,x)
```

- ❖ \$ python testbasic.py

time for one LU solve = 34.15 ms

1 or 2 magnitude difference

PyAMG

Algebraic Multigrid Solvers in Python

<http://www.pyamg.org>

Nathan Bell, Nvidia

Luke Olson, University of Illinois

Jacob Schroder, Lawrence Livermore National Lab

Thanks To NSF